

# Delhivery - Feature Engineering

## About Delhivery

Delhivery is the largest and fastest-growing fully integrated player in India by revenue in Fiscal 2021. They aim to build the operating system for commerce, through a combination of world-class infrastructure, logistics operations of the highest quality, and cutting-edge engineering and technology capabilities.

The Data team builds intelligence and capabilities using this data that helps them to widen the gap between the quality, efficiency, and profitability of their business versus their competitors.

## Business Problem

Delhivery aims to establish itself as the premier player in the logistics industry. This case study is of paramount importance as it aligns with the company's core objectives and operational excellence. It provides a practical framework for understanding and processing data, which is integral to their operations. By leveraging data engineering pipelines and data analysis techniques, Delhivery can achieve several critical goals.

First, it allows them to ensure data integrity and quality by addressing missing values and structuring the dataset appropriately. Second, it enables the extraction of valuable features from raw data, which can be utilized for building accurate forecasting models. Moreover, it facilitates the identification of patterns, insights, and actionable recommendations crucial for optimizing their logistics operations.

By conducting hypothesis testing and outlier detection, Delhivery can refine their processes and further enhance the quality of service they provide.

```
In [81]: # Lets import necessary libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from scipy.stats import norm

# Use this code to style the plots - globally
plt.style.use('ggplot')
# If you want to know what other styles are available, use plt.style.available

# Use this code to ignore any unnecessary filter warnings
import warnings
warnings.filterwarnings('ignore')
```

```
In [82]: # Lets import the dataset

df = pd.read_csv('delhivery_data.csv')
```

## 1. Basic data cleaning and exploration

### Problem Statement

Delhivery, India's largest and fastest-growing logistics company, seeks to leverage its extensive data resources to optimize its operations, increase efficiency, and maintain a competitive edge in the market. As the company processes vast amounts of raw data generated from various data engineering pipelines, there is a pressing need to transform this data into meaningful, actionable insights. The primary objective is to clean, sanitize, and engineer features from the raw data, enabling the data science team to develop accurate forecasting models that can support decision-making, enhance operational efficiency, and drive profitability.

The challenge lies in identifying and extracting the most impactful features from the raw data, ensuring the data's quality and relevance, and aligning the feature engineering process with the business goals of forecasting and predictive analytics


In [83]: *# Lets Look at the head of the data*

```
df.head()
```

Out[83]:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uuid	source_center	source_name
0	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
1	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
2	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
3	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
4	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)

5 rows × 24 columns



In [84]: *# Lets Look at the shape of the data*

```
df.shape
```

Out[84]: (144867, 24)

In [85]: *# What are the columns available in this dataset*

```
df.columns
```

Out[85]: Index(['data', 'trip\_creation\_time', 'route\_schedule\_uuid', 'route\_type', 'trip\_uuid', 'source\_center', 'source\_name', 'destination\_center', 'destination\_name', 'od\_start\_time', 'od\_end\_time', 'start\_scan\_to\_end\_scan', 'is\_cutoff', 'cutoff\_factor', 'cutoff\_timestamp', 'actual\_distance\_to\_destination', 'actual\_time', 'osrm\_time', 'osrm\_distance', 'factor', 'segment\_actual\_time', 'segment\_osrm\_time', 'segment\_osrm\_distance', 'segment\_factor'], dtype='object')

## Column Profiling

- **data** - tells whether the data is testing or training data
- **trip\_creation\_time** - Timestamp of trip creation
- **route\_schedule\_uuid** - Unique Id for a particular route schedule
- **route\_type** - Transportation type
- **FTL** - Full Truck Load: FTL shipments get to the destination sooner, as the truck is making no other pickups or drop-offs along the way
- **trip\_uuid** - Unique ID given to a particular trip (A trip may include different source and destination centers)
- **source\_center** - Source ID of trip origin
- **source\_name** - Source Name of trip origin
- **destination\_cente** - Destination ID
- **destination\_name** - Destination Name
- **od\_start\_time** - Trip start time
- **od\_end\_time** - Trip end time
- **start\_scan\_to\_end\_scan** - Time taken to deliver from source to destination
- **is\_cutoff** - Unknown field
- **cutoff\_factor** - Unknown field
- **cutoff\_timestamp** - Unknown field
- **actual\_distance\_to\_destination** - Distance in Kms between source and destination warehouse

- **actual\_time** - Actual time taken to complete the delivery (Cumulative)
- **osrm\_time** - An open-source routing engine time calculator which computes the shortest path between points in a given map (Includes usual traffic, distance through major and minor roads) and gives the time (Cumulative)
- **osrm\_distance** - An open-source routing engine which computes the shortest path between points in a given map (Includes usual traffic, distance through major and minor roads) (Cumulative)
- **factor** - Unknown field
- **segment\_actual\_time** - This is a segment time. Time taken by the subset of the package delivery
- **segment\_osrm\_time** - This is the OSRM segment time. Time taken by the subset of the package delivery
- **segment\_osrm\_distance** - This is the OSRM distance. Distance covered by subset of the package delivery
- **segment\_factor** - Unknown field

## Info of the dataframe

In [86]: *# Lets look into each column to get a better understanding of how the data looks like*

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   data                                  144867 non-null  object
1   trip_creation_time                    144867 non-null  object
2   route_schedule_uuid                  144867 non-null  object
3   route_type                           144867 non-null  object
4   trip_uuid                            144867 non-null  object
5   source_center                        144867 non-null  object
6   source_name                          144574 non-null  object
7   destination_center                   144867 non-null  object
8   destination_name                     144606 non-null  object
9   od_start_time                        144867 non-null  object
10  od_end_time                          144867 non-null  object
11  start_scan_to_end_scan                144867 non-null  float64
12  is_cutoff                            144867 non-null  bool
13  cutoff_factor                        144867 non-null  int64
14  cutoff_timestamp                     144867 non-null  object
15  actual_distance_to_destination        144867 non-null  float64
16  actual_time                          144867 non-null  float64
17  osrm_time                            144867 non-null  float64
18  osrm_distance                        144867 non-null  float64
19  factor                               144867 non-null  float64
20  segment_actual_time                  144867 non-null  float64
21  segment_osrm_time                    144867 non-null  float64
22  segment_osrm_distance                144867 non-null  float64
23  segment_factor                       144867 non-null  float64
dtypes: bool(1), float64(10), int64(1), object(12)
memory usage: 25.6+ MB
```

## Datatypes

In [87]: *# Lets check if the datatypes are correctly assigned or not*

```
# From the info above we can see that trip_creation_time, od_start_time, od_end_time and cutoff_time
# Convert those to datetime format

df['trip_creation_time'] = pd.to_datetime(df['trip_creation_time'])
df['od_start_time'] = pd.to_datetime(df['od_start_time'])
df['od_end_time'] = pd.to_datetime(df['od_end_time'])
df['cutoff_timestamp'] = pd.to_datetime(df['cutoff_timestamp'])
```

```
In [88]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144867 entries, 0 to 144866
Data columns (total 24 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   data                                    144867 non-null  object
 1   trip_creation_time                     144867 non-null  datetime64[ns]
 2   route_schedule_uuid                   144867 non-null  object
 3   route_type                             144867 non-null  object
 4   trip_uuid                             144867 non-null  object
 5   source_center                         144867 non-null  object
 6   source_name                           144574 non-null  object
 7   destination_center                    144867 non-null  object
 8   destination_name                      144606 non-null  object
 9   od_start_time                         144867 non-null  datetime64[ns]
10  od_end_time                           144867 non-null  datetime64[ns]
11  start_scan_to_end_scan                 144867 non-null  float64
12  is_cutoff                             144867 non-null  bool
13  cutoff_factor                         144867 non-null  int64
14  cutoff_timestamp                      144867 non-null  datetime64[ns]
15  actual_distance_to_destination         144867 non-null  float64
16  actual_time                           144867 non-null  float64
17  osrm_time                             144867 non-null  float64
18  osrm_distance                         144867 non-null  float64
19  factor                                144867 non-null  float64
20  segment_actual_time                   144867 non-null  float64
21  segment_osrm_time                     144867 non-null  float64
22  segment_osrm_distance                 144867 non-null  float64
23  segment_factor                        144867 non-null  float64
dtypes: bool(1), datetime64[ns](4), float64(10), int64(1), object(8)
memory usage: 25.6+ MB
```

## Null values

```
In [89]: # Lets see if there are any null values
```

```
df.isnull().sum()
```

```
Out[89]: data                                0
trip_creation_time                          0
route_schedule_uuid                         0
route_type                                 0
trip_uuid                                  0
source_center                              0
source_name                               293
destination_center                         0
destination_name                          261
od_start_time                             0
od_end_time                               0
start_scan_to_end_scan                    0
is_cutoff                                 0
cutoff_factor                             0
cutoff_timestamp                          0
actual_distance_to_destination             0
actual_time                               0
osrm_time                                 0
osrm_distance                             0
factor                                    0
segment_actual_time                       0
segment_osrm_time                         0
segment_osrm_distance                     0
segment_factor                           0
dtype: int64
```

We can see some null values in source\_name and destination\_center

```
In [90]: # Lets see the % of null values
```

```
df.isna().sum()/len(df) * 100
```

```
Out[90]: data                                0.000000
trip_creation_time                          0.000000
route_schedule_uuid                         0.000000
route_type                                  0.000000
trip_uuid                                   0.000000
source_center                              0.000000
source_name                                0.202254
destination_center                         0.000000
destination_name                           0.180165
od_start_time                              0.000000
od_end_time                                0.000000
start_scan_to_end_scan                     0.000000
is_cutoff                                  0.000000
cutoff_factor                              0.000000
cutoff_timestamp                           0.000000
actual_distance_to_destination              0.000000
actual_time                                0.000000
osrm_time                                  0.000000
osrm_distance                              0.000000
factor                                     0.000000
segment_actual_time                        0.000000
segment_osrm_time                          0.000000
segment_osrm_distance                      0.000000
segment_factor                             0.000000
dtype: float64
```

There are only .2% and .1% null values for source\_name and destination\_name respectively

```
In [91]: # Lets remove the null values
```

```
df.dropna(inplace=True)
```

## Merging Rows

```
In [92]: # Lets create a unique identifier - segment_key
```

```
df['segment_key'] = df['trip_uuid'] + '_' + df['source_center'] + '_' + df['destination_center']
```

```
In [93]: df['segment_key'].head()
```

```
Out[93]: 0    trip-153741093647649320_IND388121AAA_IND388620AAB
1    trip-153741093647649320_IND388121AAA_IND388620AAB
2    trip-153741093647649320_IND388121AAA_IND388620AAB
3    trip-153741093647649320_IND388121AAA_IND388620AAB
4    trip-153741093647649320_IND388121AAA_IND388620AAB
Name: segment_key, dtype: object
```

```
In [94]: # Now merge the rows in columns segment_actual_time, segment_osrm_distance, segment_osrm_time based
```

```
df['segment_actual_time_sum'] = df.groupby('segment_key')['segment_actual_time'].cumsum()
df['segment_osrm_distance_sum'] = df.groupby('segment_key')['segment_osrm_distance'].cumsum()
df['segment_osrm_time_sum'] = df.groupby('segment_key')['segment_osrm_time'].cumsum()
```

In [95]: *# Define aggregation rules for segment-level aggregation*

```
create_segment_dict = {
    'trip_creation_time': 'first',          # Keep the first creation time
    'route_schedule_uuid': 'first',        # Keep the first schedule
    'route_type': 'first',                 # Keep the first route type
    'od_start_time': 'first',              # Keep the start time of the segment
    'od_end_time': 'last',                 # Keep the end time of the segment
    'start_scan_to_end_scan': 'sum',       # Sum the values across rows for total duration
    'actual_distance_to_destination': 'sum', # Sum distances
    'actual_time': 'sum',                  # Sum actual times
    'osrm_time': 'sum',                    # Sum OSRM times
    'osrm_distance': 'sum',                # Sum OSRM distances
    'factor': 'mean',                      # Take mean of factors
    'segment_actual_time_sum': 'last',     # Get the last cumulative value
    'segment_osrm_distance_sum': 'last',   # Get the last cumulative value
    'segment_osrm_time_sum': 'last'       # Get the last cumulative value
}
```

In [96]: *# Aggregating at the segment level using segment\_key*

```
df_segment = df.groupby('segment_key').agg(create_segment_dict).reset_index()
```

In [97]: *# Sorting by segment\_key and then by od\_end\_time to maintain order within segments*

```
df_segment = df_segment.sort_values(by=['segment_key', 'od_end_time']).reset_index(drop=True)
```

In [98]: *# Display the resulting DataFrame*

```
df_segment.head()
```

Out[98]:

		segment_key	trip_creation_time	route_schedule_uuid	route_type	od_start_time
0	153671041653548748_IND209304AAA_IND000000ACB	trip-153671041653548748_IND209304AAA_IND000000ACB	2018-09-12 00:00:16.535741	thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6...	FTL	2018-09-1 16:39:46.85846
1	153671041653548748_IND462022AAA_IND209304AAA	trip-153671041653548748_IND462022AAA_IND209304AAA	2018-09-12 00:00:16.535741	thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6...	FTL	2018-09-1 00:00:16.53574
2	153671042288605164_IND561203AAB_IND562101AAA	trip-153671042288605164_IND561203AAB_IND562101AAA	2018-09-12 00:00:22.886430	thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0...	Carting	2018-09-1 02:03:09.65559
3	153671042288605164_IND572101AAA_IND561203AAB	trip-153671042288605164_IND572101AAA_IND561203AAB	2018-09-12 00:00:22.886430	thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0...	Carting	2018-09-1 00:00:22.88643
4	153671043369099517_IND000000ACB_IND160002AAC	trip-153671043369099517_IND000000ACB_IND160002AAC	2018-09-12 00:00:33.691250	thanos::sroute:de5e208e-7641-45e6-8100-4d9fb1e...	FTL	2018-09-1 03:40:17.10673

In [99]: df\_segment.shape

Out[99]: (26222, 15)

## Build some features to prepare the data for actual analysis

In [100]: *# Lets calculate `od\_time\_diff\_hour` by finding the time difference between `od\_start\_time` and `od\_end\_time`*

```
# Calculate the time difference in hours
df['od_time_diff_hour'] = (df['od_end_time'] - df['od_start_time']).dt.total_seconds() / 3600
```

In [101]: *# Drop `od\_start\_time` and `od\_end\_time`*

```
df = df.drop(columns=['od_start_time', 'od_end_time'])
```

```

In [102]: # Split and extract features from `destination_name` - City, Place, Code and State

df[['dest_city', 'dest_place', 'dest_code', 'dest_state']] = df['destination_name'].str.extract(r'(\w+)(\w+)(\w+)(\w+)')

In [103]: # Split and extract features from `source_name` - City, Place, Code and State

df[['source_city', 'source_place', 'source_code', 'source_state']] = df['source_name'].str.extract(r'(\w+)(\w+)(\w+)(\w+)')

In [104]: # Extract datetime features from `trip_creation_time`

df['trip_creation_year'] = df['trip_creation_time'].dt.year
df['trip_creation_month'] = df['trip_creation_time'].dt.month
df['trip_creation_day'] = df['trip_creation_time'].dt.day
df['trip_creation_hour'] = df['trip_creation_time'].dt.hour
df['trip_creation_minute'] = df['trip_creation_time'].dt.minute
df['trip_creation_weekday'] = df['trip_creation_time'].dt.weekday # Monday=0, Sunday=6

In [105]: # Drop `trip_creation_time` feature extraction

# df = df.drop(columns=['trip_creation_time'])

In [106]: # Display the resulting DataFrame
df.head()

```

Out[106]:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uuid	source_center	source_name
0	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
1	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
2	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
3	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
4	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)

5 rows × 41 columns

# In-depth analysis

## Grouping and Aggregating at Trip-level

In [107]: # Define aggregation rules for trip-level summary

```
create_trip_dict = {
    'trip_creation_time': 'first',          # Keep first creation time
    'route_schedule_uuid': 'first',        # Keep first schedule
    'route_type': 'first',                 # Keep first route type
    'od_time_diff_hour': 'sum',            # Sum up total od_time_diff
    'actual_distance_to_destination': 'sum', # Total distance to destination
    'actual_time': 'sum',                  # Total actual time
    'osrm_time': 'sum',                    # Total OSRM time
    'osrm_distance': 'sum',                # Total OSRM distance
    'factor': 'mean'                       # Mean of factor across segments
}
```

In [108]: # Group by trip\_uuid and apply aggregation

```
df_trip = df.groupby('trip_uuid').agg(create_trip_dict).reset_index()
```

In [109]: df\_trip

Out[109]:

	trip_uuid	trip_creation_time	route_schedule_uuid	route_type	od_time_diff_hour	actual_distance_to_destination
0	trip-153671041653548748	2018-09-12 00:00:16.535741	thanos::sroute:d7c989ba-a29b-4a0b-b2f4-288cdc6...	FTL	728.008209	886
1	trip-153671042288605164	2018-09-12 00:00:22.886430	thanos::sroute:3a1b0ab2-bb0b-4c53-8c59-eb2a2c0...	Carting	15.219568	2
2	trip-153671043369099517	2018-09-12 00:00:33.691250	thanos::sroute:de5e208e-7641-45e6-8100-4d9fb1e...	FTL	4144.906395	6816
3	trip-153671046011330457	2018-09-12 00:01:00.113710	thanos::sroute:f0176492-a679-4597-8332-bbd1c7f...	Carting	3.349831	4
4	trip-153671052974046625	2018-09-12 00:02:09.740725	thanos::sroute:d9f07b12-65e0-4f3b-bec8-df06134...	FTL	26.478517	2
...	...	...	...	...	...	...
14782	trip-153861095625827784	2018-10-03 23:55:56.258533	thanos::sroute:8a120994-f577-4491-9e4b-b7e4a14...	Carting	14.655464	1
14783	trip-153861104386292051	2018-10-03 23:57:23.863155	thanos::sroute:b30e1ec3-3bfa-4bd2-a7fb-3b75769...	Carting	2.019684	4
14784	trip-153861106442901555	2018-10-03 23:57:44.429324	thanos::sroute:5609c268-e436-4e0a-8180-3db4a74...	Carting	21.105993	9
14785	trip-153861115439069069	2018-10-03 23:59:14.390954	thanos::sroute:c5f2ba2c-8486-4940-8af6-d1d2a6a...	Carting	22.006709	3
14786	trip-153861118270144424	2018-10-03 23:59:42.701692	thanos::sroute:412fea14-6d1f-4222-8a5f-a517042...	FTL	11.813586	1

14787 rows × 10 columns





## Outlier Detection & Treatment

```
In [110]: # We can detect outliers using IQR method

# Select numeric columns

numeric_cols = df_trip.select_dtypes(include=[np.number]).columns
```

```
In [111]: # Calculate Q1 and Q3

Q1 = df_trip[numeric_cols].quantile(0.25)
Q3 = df_trip[numeric_cols].quantile(0.75)
IQR = Q3 - Q1
```

```
In [112]: # Detect outliers


outliers = ((df_trip[numeric_cols] < (Q1 - 1.5 * IQR)) | (df_trip[numeric_cols] > (Q3 + 1.5 * IQR)))
```

```
In [113]: # Lets visualize the outliers using a Boxplot
```

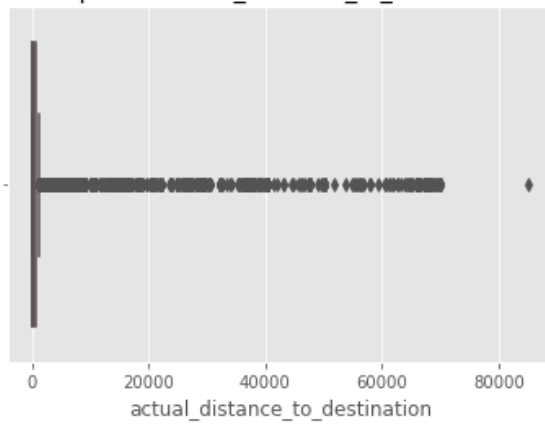
```
# Plotting boxplots for each numeric column
```

```
for col in numeric_cols:  
    plt.figure(figsize=(6, 4))  
    sns.boxplot(x=df_trip[col])  
    plt.title(f'Boxplot of {col}')  
    plt.show()
```


Boxplot of od\_time\_diff\_hour

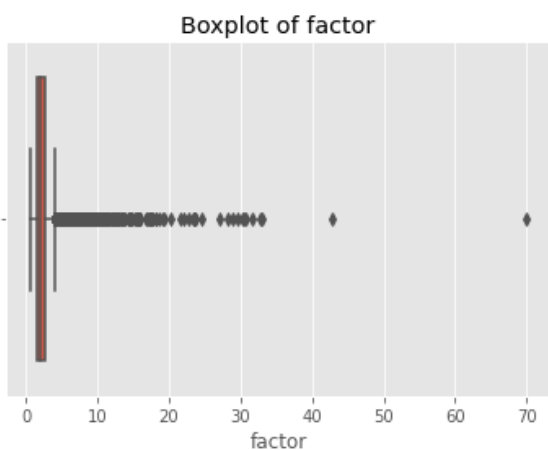

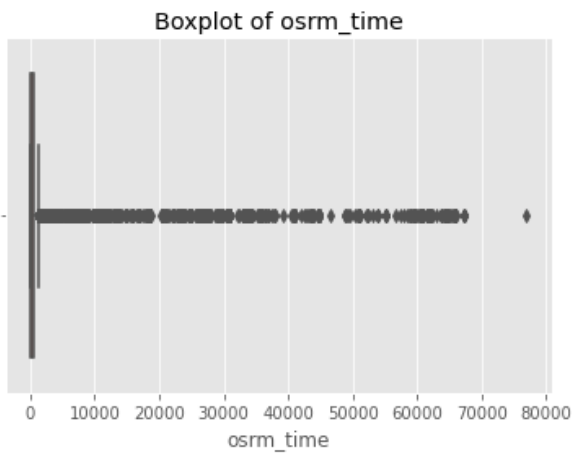


Boxplot of actual\_distance\_to\_destination



Boxplot of actual\_time





```
In [114]: # We can see a lot of outliers, Lets handle them using IQR method

# Capping outliers to 1.5 * IQR Limits
for col in numeric_cols:
    lower_bound = Q1[col] - 1.5 * IQR[col]
    upper_bound = Q3[col] + 1.5 * IQR[col]
    df_trip[col] = np.where(df_trip[col] < lower_bound, lower_bound, df_trip[col])
    df_trip[col] = np.where(df_trip[col] > upper_bound, upper_bound, df_trip[col])
```

## One-Hot Encoding of Categorical Features

```
In [115]: # Identify categorical columns for encoding

categorical_cols = df_trip.select_dtypes(include=['object']).columns
```

```
In [116]: # Apply one-hot encoding

df_trip = pd.get_dummies(df_trip, columns=categorical_cols, drop_first=True)
```

## Normalize/Standardize Numerical Features

```
In [118]: from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

```
In [119]: # We have two methods for scaling - MinMaxScaler() and StandardScaler()

# Lets use MinMaxScaler

scaler = MinMaxScaler()      # Normalization
# scaler = StandardScaler()  # Standardization
```

```
In [120]: # Apply scaling to numerical columns

df_trip[numeric_cols] = scaler.fit_transform(df_trip[numeric_cols])
```

## Hypothesis Testing

### Aggregating Values by trip\_uuid

```
In [138]: # Aggregating necessary columns at the trip level

df_trip_agg = df.groupby('trip_uuid').agg({
    'actual_time': 'sum',
    'osrm_time': 'sum',
    'segment_actual_time': 'sum',
    'segment_osrm_time': 'sum',
    'osrm_distance': 'sum',
    'segment_osrm_distance': 'sum'
}).reset_index()
```

```
In [139]: # Lets handle the outliers

numeric_cols = df_trip_agg.select_dtypes(include=[np.number]).columns

Q1 = df_trip_agg[numeric_cols].quantile(0.25)
Q3 = df_trip_agg[numeric_cols].quantile(0.75)
IQR = Q3 - Q1

outliers = ((df_trip_agg[numeric_cols] < (Q1 - 1.5 * IQR)) | (df_trip_agg[numeric_cols] > (Q3 + 1.5 * IQR)))

for col in numeric_cols:
    lower_bound = Q1[col] - 1.5 * IQR[col]
    upper_bound = Q3[col] + 1.5 * IQR[col]
    df_trip_agg[col] = np.where(df_trip_agg[col] < lower_bound, lower_bound, df_trip_agg[col])
    df_trip_agg[col] = np.where(df_trip_agg[col] > upper_bound, upper_bound, df_trip_agg[col])
```

## Visual Analysis

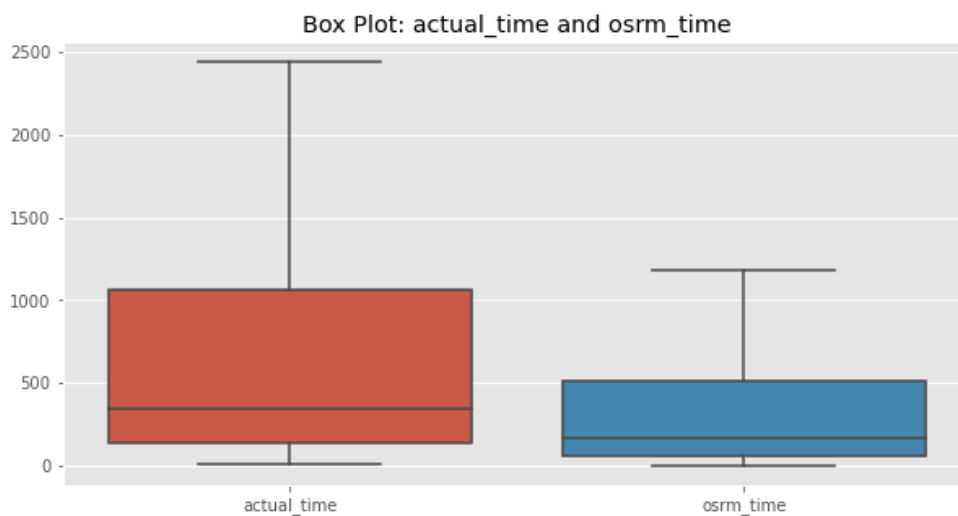

```
In [140]: # Define pairs to compare


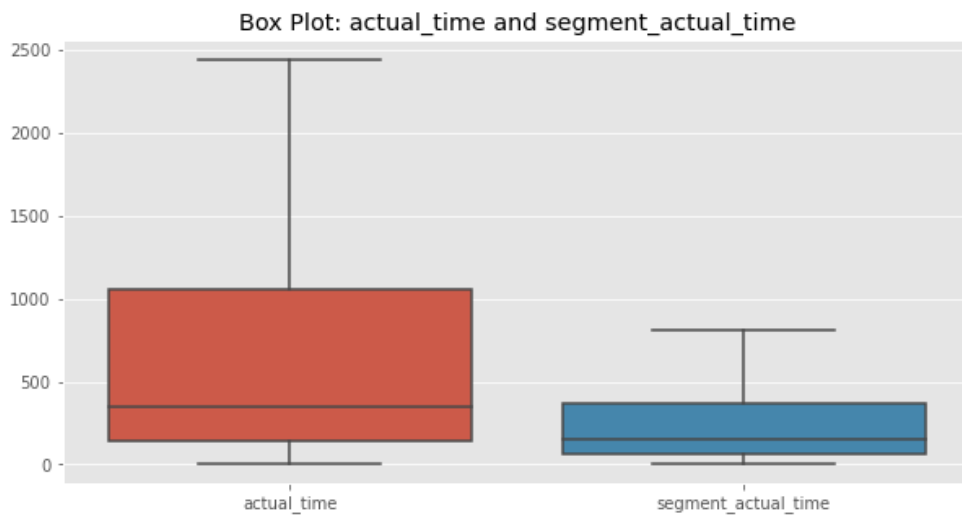

comparison_pairs = [
    ('actual_time', 'osrm_time'),
    ('actual_time', 'segment_actual_time'),
    ('osrm_distance', 'segment_osrm_distance'),
    ('osrm_time', 'segment_osrm_time')
]
```




```
In [141]: # Plotting scatter and box plots for each pair
```

```
for x, y in comparison_pairs:
    plt.figure(figsize=(10, 5))
    sns.scatterplot(x=df_trip_agg[x], y=df_trip_agg[y])
    plt.title(f'Scatter Plot: {x} vs {y}')
    plt.xlabel(x)
    plt.ylabel(y)
    plt.show()

    plt.figure(figsize=(10, 5))
    sns.boxplot(data=df_trip_agg[[x, y]])
    plt.title(f'Box Plot: {x} and {y}')
    plt.show()
```







## Hypothesis Testing

```
In [142]: from scipy.stats import ttest_rel
```

In [143]: *# Perform paired t-tests for each comparison pair*

```
for x, y in comparison_pairs:
    t_stat, p_value = ttest_rel(df_trip_agg[x], df_trip_agg[y])
    print(f'Hypothesis Test for {x} vs {y}:')
    print(f'T-statistic = {t_stat:.3f}, P-value = {p_value:.3f}\n')
```

Hypothesis Test for actual\_time vs osrm\_time:  
T-statistic = 104.254, P-value = 0.000

Hypothesis Test for actual\_time vs segment\_actual\_time:  
T-statistic = 99.494, P-value = 0.000

Hypothesis Test for osrm\_distance vs segment\_osrm\_distance:  
T-statistic = 97.276, P-value = 0.000

Hypothesis Test for osrm\_time vs segment\_osrm\_time:  
T-statistic = 99.371, P-value = 0.000

## Business Insights & Recommendations


### Identify Sources of Most Orders

In [144]: `df.head()`

Out[144]:

	data	trip_creation_time	route_schedule_uuid	route_type	trip_uuid	source_center	source_name
0	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
1	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
2	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
3	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)
4	training	2018-09-20 02:35:36.476840	thanos::sroute:eb7bfc78- b351-4c0e-a951- fa3d5c3...	Carting	153741093647649320	IND388121AAA	Anand_VUNagar_DC (Gujarat)

5 rows × 42 columns



In [145]: *# Most frequent source and destination states or centers*

```
most_orders_source = df['source_name'].value_counts().head(10)
most_orders_destination = df['destination_name'].value_counts().head(10)
```



In [146]: *# Display results*

```
print("Top 10 Source Centers:\n", most_orders_source)
print("\n")
print("Top 10 Destination Centers:\n", most_orders_destination)
```

Top 10 Source Centers:

Gurgaon_Bilaspur_HB (Haryana)	23267
Bangalore_Nelmngla_H (Karnataka)	9975
Bhiwandi_Mankoli_HB (Maharashtra)	9088
Pune_Tathawde_H (Maharashtra)	4061
Hyderabad_Shamshbd_H (Telangana)	3340
Kolkata_Dankuni_HB (West Bengal)	2612
Chandigarh_Mehmdpur_H (Punjab)	2450
Surat_HUB (Gujarat)	2189
Delhi_Airport_H (Delhi)	1997
Bengaluru_Bomsndra_HB (Karnataka)	1958

Name: source\_name, dtype: int64

Top 10 Destination Centers:

Gurgaon_Bilaspur_HB (Haryana)	15192
Bangalore_Nelmngla_H (Karnataka)	11019
Bhiwandi_Mankoli_HB (Maharashtra)	5492
Hyderabad_Shamshbd_H (Telangana)	5142
Kolkata_Dankuni_HB (West Bengal)	4892
Delhi_Airport_H (Delhi)	3761
Pune_Tathawde_H (Maharashtra)	3695
Chandigarh_Mehmdpur_H (Punjab)	2874
Sonipat_Kundli_H (Haryana)	2796
Bhubaneshwar_Hub (Orissa)	2524

Name: destination\_name, dtype: int64

## Analyze Busiest Corridors, Average Distance, and Time Taken

In [147]: *# Create a corridor identifier*

```
df['corridor'] = df['source_name'] + " to " + df['destination_name']
```

In [148]: *# Find the busiest corridors*

```
busiest_corridors = df['corridor'].value_counts().head(10)
print("Top 10 Busiest Corridors:\n", busiest_corridors)
```

Top 10 Busiest Corridors:

Gurgaon_Bilaspur_HB (Haryana) to Bangalore_Nelmngla_H (Karnataka)	4976
Bangalore_Nelmngla_H (Karnataka) to Gurgaon_Bilaspur_HB (Haryana)	3316
Gurgaon_Bilaspur_HB (Haryana) to Kolkata_Dankuni_HB (West Bengal)	2862
Gurgaon_Bilaspur_HB (Haryana) to Hyderabad_Shamshbd_H (Telangana)	1639
Gurgaon_Bilaspur_HB (Haryana) to Bhiwandi_Mankoli_HB (Maharashtra)	1617
Bhiwandi_Mankoli_HB (Maharashtra) to Gurgaon_Bilaspur_HB (Haryana)	1269
Guwahati_Hub (Assam) to Delhi_Airport_H (Delhi)	1137
Bhiwandi_Mankoli_HB (Maharashtra) to Bangalore_Nelmngla_H (Karnataka)	1131
Gurgaon_Bilaspur_HB (Haryana) to Pune_Tathawde_H (Maharashtra)	1120
Gurgaon_Bilaspur_HB (Haryana) to MAA_Poonamallee_HB (Tamil Nadu)	1015

Name: corridor, dtype: int64

## Average Distance and Time Taken for Each Corridor

In [149]: # Group by corridor to get average distance and time taken

```
corridor_stats = df.groupby('corridor').agg({
    'actual_distance_to_destination': 'mean',
    'actual_time': 'mean'
}).rename(columns={
    'actual_distance_to_destination': 'avg_distance',
    'actual_time': 'avg_time'
}).reset_index()
```

In [150]: # Display corridor stats for busiest corridors

```
print("Average Distance and Time for Busiest Corridors:\n", corridor_stats.head(10))
```

Average Distance and Time for Busiest Corridors:

	corridor	avg_distance	avg_time
0	AMD_Memnagar (Gujarat) to Ahmedabad_East_H_1 (...)	13.166738	30.437500
1	AMD_Rakhial (Gujarat) to Ahmedabad_East_H_1 (G...	11.704146	43.603175
2	Abohar_DC (Punjab) to Malout_DC (Punjab)	22.113336	19.000000
3	Abohar_DC (Punjab) to Muktsar_DPC (Punjab)	37.613637	76.000000
4	Achrol_BgwriDPP_D (Rajasthan) to Jaipur_Hub (R...	31.059072	61.000000
5	Addanki_Oilmilrd_D (Andhra Pradesh) to Ongole_...	29.661637	174.000000
6	Adoor_Town_D (Kerala) to Kollam_Central_H_1 (K...	22.179028	93.594595
7	Agra_Central_D_3 (Uttar Pradesh) to Kirauli_Ac...	17.253622	34.350000
8	Agra_Idgah_L (Uttar Pradesh) to Gurgaon_Bilasp...	97.921666	171.437500
9	Agra_Idgah_P (Uttar Pradesh) to Delhi_Airport_...	93.981262	220.512195

## Business Insights and Recommendations

From the above analyses, here are some patterns and actionable insights:

**Source and Destination Patterns:** We have identified the top Source states as *Haryana, Karnataka, Maharashtra, Telangana, West Bengal, Punjab, Gujarat, Delhi* and the top Destination states as *Haryana, Karnataka, Maharashtra, Telangana, West Bengal, Delhi, Punjab, Orissa* Recommendation: Scale infrastructure and improve resources (e.g., warehouses, staff) in these high-demand locations to meet growing order volumes and minimize delays

**Busiest Corridors:** Some of the busiest corridors are Gurgaon to Bangalore, Bangalore to Gurgaon, Gurgaon to Kolkata etc Recommendation: Increase logistics resources on popular routes to handle high volumes efficiently. Additionally, for corridors with consistently high travel times, explore alternative routes or transit methods to reduce delivery times

**Distance and Time Analysis:** Corridors with long average distances but high volume might require optimization or faster modes of transport Recommendation: Introduce express shipping options or partnerships on long corridors with high traffic, ensuring faster deliveries and increased customer satisfaction

In [ ]: